

Static vs Contextual Embeddings for Lexical Ambiguity: A Word-in-Context Evaluation Method and Results Appendix

Choudhary Prashant Santosh (1910474)

Rahul Khunt (1911272)

September 2026

Contents

1 Research question and scope	1
2 Dataset provenance and audit	2
3 Representations	2
3.1 Static baselines	2
3.2 Contextual target	2
4 Selection and evaluation protocol	2
5 Primary validation results	3
6 Secondary BERT layer analysis	5
7 Error analysis	6
8 Reproducibility record	6
9 Limitations	6
10 Author contributions and use of AI tools	6
11 References	7

Project repository: <https://github.com/Mr-Dark-debug/static-contextual-ambiguity>

1 Research question and scope

This project asks whether a context-sensitive target representation distinguishes same-sense from different-sense uses more effectively than static alternatives. It compares type-level GloVe, a lexical GloVe context mean, and an unfine-tuned BERT-base target representation. It is a controlled representation probe, not a claim about fine-tuned task state of the art.

2 Dataset provenance and audit

The source is the official SuperGLUE v2 WiC archive at https://dl.fbaipublicfiles.com/glu_e/superglue/data/v2/WiC.zip. Its SHA-256 is ee7e67f4ae9eafbf533780faa198e62167f3cda54256cdf261877be3c0e90900. The release is distributed under CC BY-NC 4.0 by the WiC authors.

Table 1: Validated release counts and span diagnostics.

Split	Rows	Positive	Negative	Unlabeled	Lemma mismatch	Repeated surface
Train	5,428	2,714	2,714	0	1,852	93
Validation	638	319	319	0	251	1
Test	1,400	0	0	1,400	521	20

All target spans are non-empty half-open character intervals. IDs are unique within splits. Four symmetric duplicate pairs occur within training, none within validation, and no task pair crosses train and validation. The label-hidden test split is audited but not evaluated.

3 Representations

3.1 Static baselines

GloVe uses the 6B-token, uncased, 300-dimensional release. Archive SHA-256: 617afb2fe6cb0d085c235baf7a465b96f4112bd7f7ccb2b2cbd649fed9cbcf2fb. The target diagnostic looks up the supplied lemma on both sides, making every available paired cosine exactly 1 by construction. Eighteen training lemmas are OOV; missing target scores are replaced by the median finite training score.

Context vectors average in-vocabulary regex word tokens after excluding only the token span that overlaps the marked target. Candidate windows are ± 2 , ± 5 , ± 10 , and the sentence. If a context side has no usable token, it falls back to that side’s target lemma; this occurs once per window in training and zero times in validation. No evaluated context pair remains missing.

3.2 Contextual target

The encoder is google-bert/bert-base-uncased, revision f584a14f45f850d60d0036d54f3e4dfdb3d52075. It remains in evaluation mode; there is no fine-tuning. A fast tokenizer emits character offsets. Every WordPiece satisfying `token_end > target_start` and `token_start < target_end` is selected, ignoring zero-width special tokens, and selected WordPieces are averaged. Layers 1-12 and mean-last-four are cached from one forward pass.

4 Selection and evaluation protocol

The balanced training labels are stratified into 4,342 tune and 1,086 internal holdout rows using seed 2026. Candidate thresholds are fitted on tune rows by macro F1. Candidate representations are ranked on internal-holdout macro F1, then accuracy, ROC-AUC, and declared simplicity order. Winners are frozen before validation metrics are computed. Thresholds are then refitted on all 5,428 training rows. Selection ledger hash: debfd61b068f22ea1a737f238bd90ea4fb8bf44f675907e428949b559a0bb309.

Table 2: Training-only candidate selection.

Family	Candidate	Tune macro F1	Holdout macro F1	Selected
Static context	glove_context_2	0.5840	0.5961	yes
Static context	glove_context_5	0.5883	0.5928	no
Static context	glove_context_10	0.5917	0.5863	no
Static context	glove_context_sentence	0.5930	0.5845	no
BERT	layer_12	0.6789	0.7145	no
BERT	mean_last_four	0.6891	0.7247	yes

5 Primary validation results

Predictions use score $\geq$ threshold as same-sense. Precision and recall use same-sense as the positive label. Confusion-matrix layout is TN, FP / FN, TP. Intervals are percentile intervals from 1,000 fixed-seed paired example resamples.

Table 3: Primary results on 638 balanced validation pairs.

System	Threshold	Accuracy	Accuracy 95% CI	Macro F1	Precision	Recall	ROC-AUC
GloVe target	1.0000	0.5000	[0.4639, 0.5392]	0.3333	0.5000	1.0000	0.5000
GloVe context (± 2)	0.6506	0.5549	[0.5156, 0.5956]	0.5511	0.5670	0.4639	0.5800
BERT mean last four	0.5813	0.6708	[0.6364, 0.7085]	0.6673	0.6416	0.7743	0.7163

The selected BERT representation is correct on 428 pairs; GloVe context is correct on 354. BERT alone is correct on 178 pairs, GloVe context alone on 104, both on 250, and both primary systems are wrong on 106.

Table 4: Paired primary-system differences with 95% intervals.

Contrast	Metric	Estimate	95% interval
glove_target - glove_context_2	accuracy	-0.0549	[-0.1176, 0.0094]
glove_target - glove_context_2	macro f1	-0.2178	[-0.2661, -0.1710]
glove_target - glove_context_2	roc auc	-0.0800	[-0.1273, -0.0378]
glove_target - bert_mean_last_four	accuracy	-0.1708	[-0.2195, -0.1238]
glove_target - bert_mean_last_four	macro f1	-0.3340	[-0.3739, -0.2953]
glove_target - bert_mean_last_four	roc auc	-0.2163	[-0.2588, -0.1775]
glove_context_2 - bert_mean_last_four	accuracy	-0.1160	[-0.1693, -0.0658]
glove_context_2 - bert_mean_last_four	macro f1	-0.1161	[-0.1690, -0.0664]
glove_context_2 - bert_mean_last_four	roc auc	-0.1363	[-0.1912, -0.0811]

The accuracy, macro-F1, and ROC-AUC intervals for GloVe context minus BERT all exclude zero. This supports a difference under the current paired sample and protocol; it does not provide a universal model ordering.

6 Secondary BERT layer analysis

Each layer threshold is trained on all training pairs. Validation layer results are descriptive and do not replace the frozen primary representation.

Table 5: Individual BERT target layers.

Layer	Threshold	Accuracy	Macro F1	ROC-AUC
1	0.8546	0.5705	0.5685	0.5900
2	0.8145	0.6019	0.6014	0.6187
3	0.7590	0.6113	0.6113	0.6357
4	0.7330	0.6317	0.6314	0.6575
5	0.6949	0.6458	0.6451	0.6768
6	0.6564	0.6442	0.6419	0.6853
7	0.6299	0.6489	0.6467	0.6946
8	0.6039	0.6599	0.6558	0.7036
9	0.5728	0.6661	0.6624	0.7093
10	0.5625	0.6599	0.6543	0.7197
11	0.5947	0.6536	0.6508	0.7096
12	0.4942	0.6364	0.6308	0.6971

Accuracy rises broadly from layer 1 to layer 9 and then declines through layer 12. Layer 10 has the largest individual-layer ROC-AUC. This is compatible with the train-selected mean-last-four result, but validation was not reused for selection.

7 Error analysis

System-agreement partitions are disjoint. There are 250 both-primary-correct, 178 BERT-only-correct, 104 static-only-correct, 54 all-system-wrong, and 52 rows where both primary systems are wrong but the one-class target diagnostic is right. For BERT, error margin is absolute distance from the frozen cosine threshold. Of 210 errors, 33 are within 0.025 and 53 occupy the top error-margin quartile.

BERT recovers 77.43% of same-sense pairs and 56.74% of different-sense pairs; GloVe context recovers 46.39% and 64.58%, respectively. These are descriptive class skews. Saved high-margin examples include a same-sense anatomical *ball/balls* pair that BERT separates strongly and different-sense *engrave* and *map* pairs that it places close together. Conversely, BERT links low-overlap same-sense examples such as *rust remover/paint remover*. Fine sense granularity and constructional similarity are plausible explanations, not causal findings.

8 Reproducibility record

The locked runtime is Python 3.11.15, NumPy 2.4.6, PyTorch 2.11.0+cu130, and Transformers 5.17.0. Full inference ran on an NVIDIA GeForce RTX 3050 Ti Laptop GPU. Arrays are keyed by dataset fingerprint, model revision, maximum length, seed, and cache schema. The clean cached reproduction recorded Git commit 7c6c956a173920b8aa685e206b9e4cdade1afdb6 on branch `main`.

9 Limitations

The sample covers one English benchmark; bootstrap intervals omit dataset-design and modeling uncertainty. BERT is an older frozen base encoder. The static context baseline has no syntax, weighting, subword model, or supervised classifier. A global cosine threshold compresses potentially useful geometry. Surface-mismatch groups are observational and confounded. No score is reported for the label-hidden test split.

10 Author contributions and use of AI tools

AI-assisted wording notice: OpenAI Codex drafted the wording of this section for the authors to verify before signing their declarations.

The authors selected WiC, designed the three comparison conditions, initialized the frozen GloVe and BERT checkpoints, and chose settings and thresholds using training data only. They ran and checked the audit, inference, validation analysis, tables, and figures. They did not train or fine-tune model weights. The saved selection ledger, predictions, and metrics record these choices.

The authors chose the poster’s explanations, examples, and graphs and checked the numbers against saved results. Canva helped explore the layout; the final poster was assembled in LaTeX. GLM assisted with LaTeX layout code. OpenAI Codex provided limited coding suggestions and debugging in an interactive paired-programming process. It also drafted this disclosure, revised appendix wording and the poster’s final italic sentence, and checked the PDFs. The authors remain responsible for the code, claims, figures, and final submission. No AI detector score is used as evidence of authorship.

11 References

References

- [1] J. Camacho-Collados and M. T. Pilehvar. From Word to Sense Embeddings: A Survey on Vector Representations of Meaning. *JAIR*, 2018. <https://doi.org/10.1613/jair.1.11259>.
- [2] M. T. Pilehvar and J. Camacho-Collados. WiC: The Word-in-Context Dataset for Evaluating Context-Sensitive Meaning Representations. *NAACL*, 2019. <https://doi.org/10.18653/v1/N19-1128>.
- [3] A. Wang et al. SuperGLUE: A Stickier Benchmark for General-Purpose Language Understanding Systems. *NeurIPS*, 2019. https://papers.nips.cc/paper_files/paper/2019/hash/4496bf24afe7fab6f046bf4923da8de6-Abstract.html.
- [4] J. Pennington, R. Socher, and C. D. Manning. GloVe: Global Vectors for Word Representation. *EMNLP*, 2014. <https://doi.org/10.3115/v1/D14-1162>.
- [5] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. *NAACL*, 2019. <https://doi.org/10.18653/v1/N19-1423>.
- [6] D. Loureiro et al. Analysis and Evaluation of Language Models for Word Sense Disambiguation. *Computational Linguistics*, 2021. https://doi.org/10.1162/coli_a_00405.
- [7] Y. Zhou and D. Bollegala. Learning Sense-Specific Static Embeddings using Contextualised Word Embeddings as a Proxy. *PACLIC*, 2021. <https://aclanthology.org/2021.paclic-1.53/>.
- [8] S. Trott and B. Bergen. RAW-C: Relatedness of Ambiguous Words in Context. *ACL-IJCNLP*, 2021. <https://doi.org/10.18653/v1/2021.acl-long.550>.
- [9] J. Haber and M. Poesio. Patterns of Polysemy and Homonymy in Contextualised Language Models. *Findings of EMNLP*, 2021. <https://doi.org/10.18653/v1/2021.findings-emnlp.226>.
- [10] E. Soper and J.-P. Koenig. When Polysemy Matters: Modeling Semantic Categorization with Word Embeddings. **SEM*, 2022. <https://doi.org/10.18653/v1/2022.starsem-1.10>.
- [11] J. Fodor, S. De Deyne, and S. Suzuki. The Importance of Context in the Evaluation of Word Embeddings: The Effects of Antonymy and Polysemy. *IWCS*, 2023. <https://aclanthology.org/2023.iwcs-1.16/>.
- [12] T. C. Erren and P. E. Bourne. Ten Simple Rules for a Good Poster Presentation. *PLOS Computational Biology*, 2007. <https://doi.org/10.1371/journal.pcbi.0030102>.
- [13] SuperGLUE. WiC v2 dataset archive. <https://dl.fbaipublicfiles.com/glue/superglue/data/v2/WiC.zip>.
- [14] Stanford NLP. GloVe pretrained-vector release, 6B tokens, 300 dimensions. <https://nlp.stanford.edu/projects/glove/>.
- [15] Google BERT. bert-base-uncased model card and checkpoint. <https://huggingface.co/google-bert/bert-base-uncased>.

Declaration of Academic Integrity

I,

Last name, first name: Choudhary Prashant Santosh

Student ID number: 1910474

hereby declare that I have written this thesis independently and have not used any sources or aids other than those specified. All direct or indirect content and ideas from external sources are acknowledged as references. I confirm that I have not previously submitted this thesis in part or whole for any other academic credit nor has the work been submitted to any other examination office and has not been published.

I confirm that I have only used text or other content-generating aids based on generative artificial intelligence (GenAI) in the manner permitted in writing by the examiner and that I have explicitly marked GenAI-generated passages.

When using GenAI tools, I have indicated these with their product name and listed them in full in an appendix.

I am fully responsible for any GenAI-generated passages and content I have used in my work.

Thesis title: Static vs Contextual Embeddings for Lexical Ambiguity

Programme: MSc Natural Language Processing

Examiner: _____

This is an achievement for a

☐ Bachelor's thesis

☐ Master's thesis

☒ other examination achievement for module (term paper etc.) Trends in Natural Language Processing

Place, Date

Signature

Declaration of Academic Integrity

I,

Last name, first name: Rahul Khunt

Student ID number: 1911272

hereby declare that I have written this thesis independently and have not used any sources or aids other than those specified. All direct or indirect content and ideas from external sources are acknowledged as references. I confirm that I have not previously submitted this thesis in part or whole for any other academic credit nor has the work been submitted to any other examination office and has not been published.

I confirm that I have only used text or other content-generating aids based on generative artificial intelligence (GenAI) in the manner permitted in writing by the examiner and that I have explicitly marked GenAI-generated passages.

When using GenAI tools, I have indicated these with their product name and listed them in full in an appendix.

I am fully responsible for any GenAI-generated passages and content I have used in my work.

Thesis title: Static vs Contextual Embeddings for Lexical Ambiguity

Programme: MSc Natural Language Processing

Examiner: _____

This is an achievement for a

☐ Bachelor's thesis

☐ Master's thesis

☒ other examination achievement for module (term paper etc.) Trends in Natural Language Processing

Place, Date

Signature